

Plots

September 3, 2017

1 Some plots and analysis of simulation data

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import math
%matplotlib inline
```

1.1 Simulation Parameters

```
In [2]: test_info = pd.read_csv("../data/backup/stats_mlfd/test_info.csv")
test_info
```

```
Out [2]:
```

	test	workersCount	locationsCount	sampleSize	defaultMean	\
0	mlfd_test	100	100	100	3000.0	
	hbTimeout	randomMillis	geoFactor	crashProb	delta	
0	4 seconds	200	100.0	0.001	NaN	


1.2 Prediction Evaluation of the Online-learner in the FD

```
In [3]: test_evaluations = pd.read_csv("../data/backup/stats_mlfd/testresults.csv")
test_evaluations.head()
```

```
Out [3]:
```

	meanSquaredError	rootMeanSquaredError	rSquared	meanAbsoluteError	\
0	2.891005e+06	1700.295700	-0.815136	1418.753397	
1	2.135465e+05	462.110885	0.974527	349.727503	
2	1.632067e+04	127.752381	0.996729	105.777594	
3	4.299215e+03	65.568400	0.999504	56.537756	
4	3.624345e+03	60.202532	0.999099	48.894390	
	explainedVariance	timestamp			
0	6.844245e+06	1504424870000			
1	1.041643e+07	1504424876000			
2	5.328207e+06	1504424883000			
3	8.820089e+06	1504424888000			
4	4.148113e+06	1504424895000			

```
In [4]: time = test_evaluations["timestamp"].values
mse = test_evaluations["meanSquaredError"].values
plt.plot(time, mse)
plt.ylabel('Mean Squared Error')
plt.xlabel('timestamps')
plt.title('MSE for timeout predictions with respect to time')
plt.show()
```



1.3 RTT Distributions related to geographic location

```
In [5]: rtt_data = pd.read_csv('../data/backup/stats_mlfd/rtt_data.csv')
```

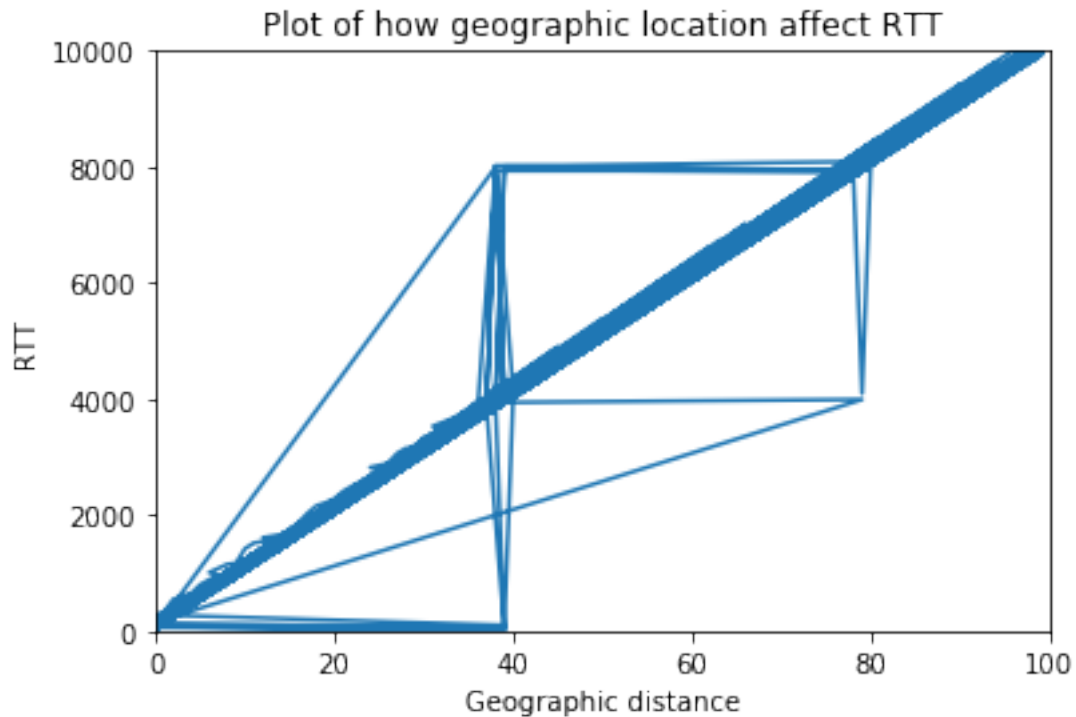
```
In [6]: rtt_data.head()
```

```
Out[6]:
```

	node	geoLocation	rtt	timestamp	mean
0	100	0	270.0	1504424862481	270.0
1	1	1	338.0	1504424862550	338.0
2	2	2	548.0	1504424862760	548.0
3	3	3	599.0	1504424862811	599.0
4	4	4	648.0	1504424862860	648.0

```
In [7]: geo = rtt_data["geoLocation"].values
rtt = rtt_data["rtt"].values
plt.plot(geo, rtt)
```

```
plt.axis([0, 100, 0, 10000])
plt.ylabel('RTT')
plt.xlabel('Geographic distance')
plt.title('Plot of how geographic location affect RTT')
plt.show()
```



```
In [8]: rtt_data_with_mean = rtt_data.groupby(['geoLocation', 'rtt'], as_index=False).mean().groupby('geoLocation')
rtt_data_with_mean.head()
```

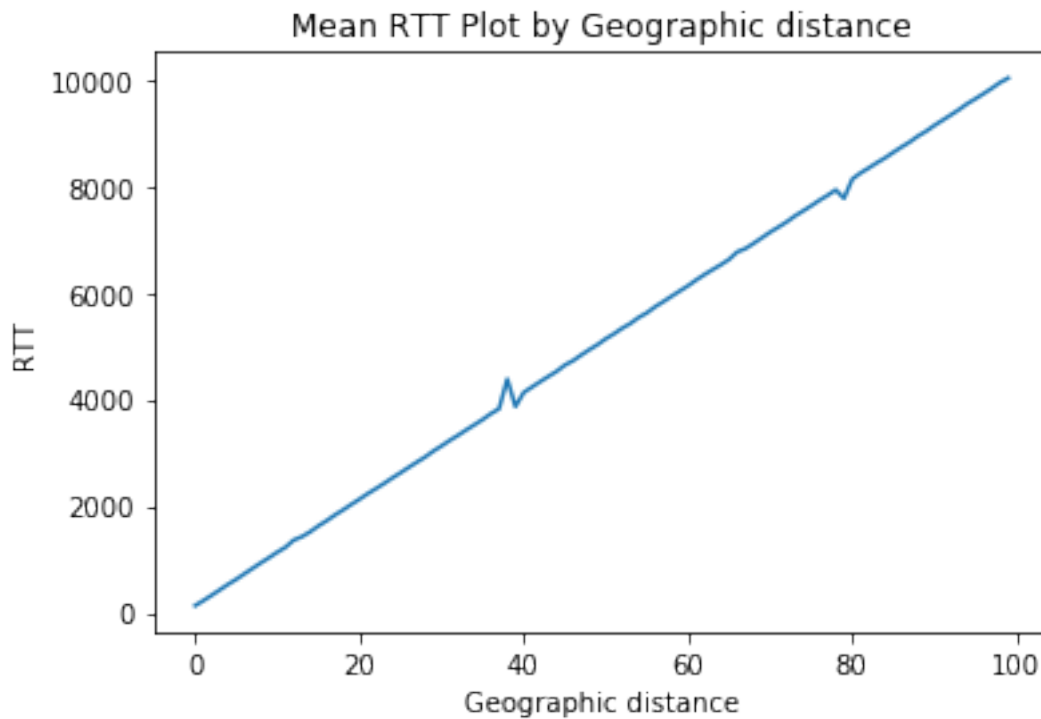
```
Out[8]: geoLocation
0      131.860000
1      226.456522
2      325.627907
3      427.609756
4      530.630435
Name: rtt, dtype: float64
```

```
In [9]: rtt_data_with_std = rtt_data.groupby('geoLocation')[['rtt']].std()
rtt_data_with_std.head()
```

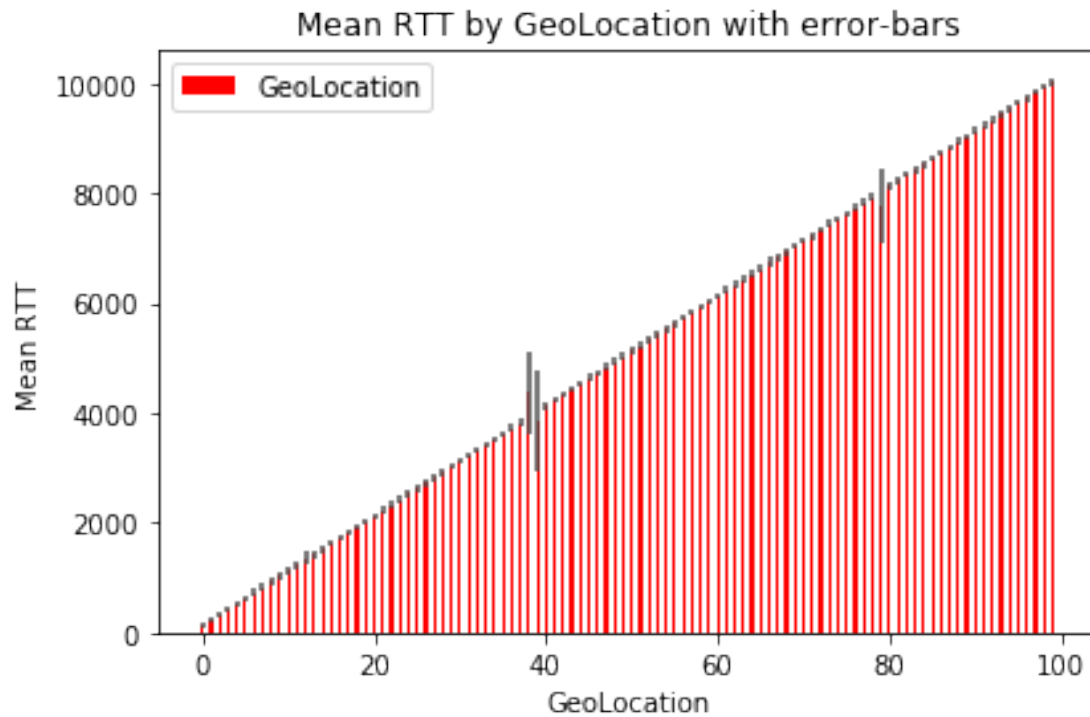
```
Out[9]:
          rtt
geoLocation
0      58.987944
1      57.103708
```

2	61.323999
3	59.254465
4	60.819616


```
In [10]: plt.plot(rtt_data_with_mean)
          #plt.axis([0, 100, 0, 10000])
          plt.ylabel('RTT')
          plt.xlabel('Geographic distance')
          plt.title('Mean RTT Plot by Geographic distance')
          plt.show()
```



```
In [11]: rtt = rtt_data_with_mean.values
          std = rtt_data_with_std.values
          l = len(rtt)
          width = 0.35
          error_config = {'ecolor': '0.3'}
          plt.bar(range(len(rtt)), rtt, width=width, color='r', yerr=std, error_kw=error_config,
                  plt.xlabel('GeoLocation')
                  plt.ylabel('Mean RTT')
                  plt.title('Mean RTT by GeoLocation with error-bars')
                  plt.legend()
                  plt.tight_layout()
```

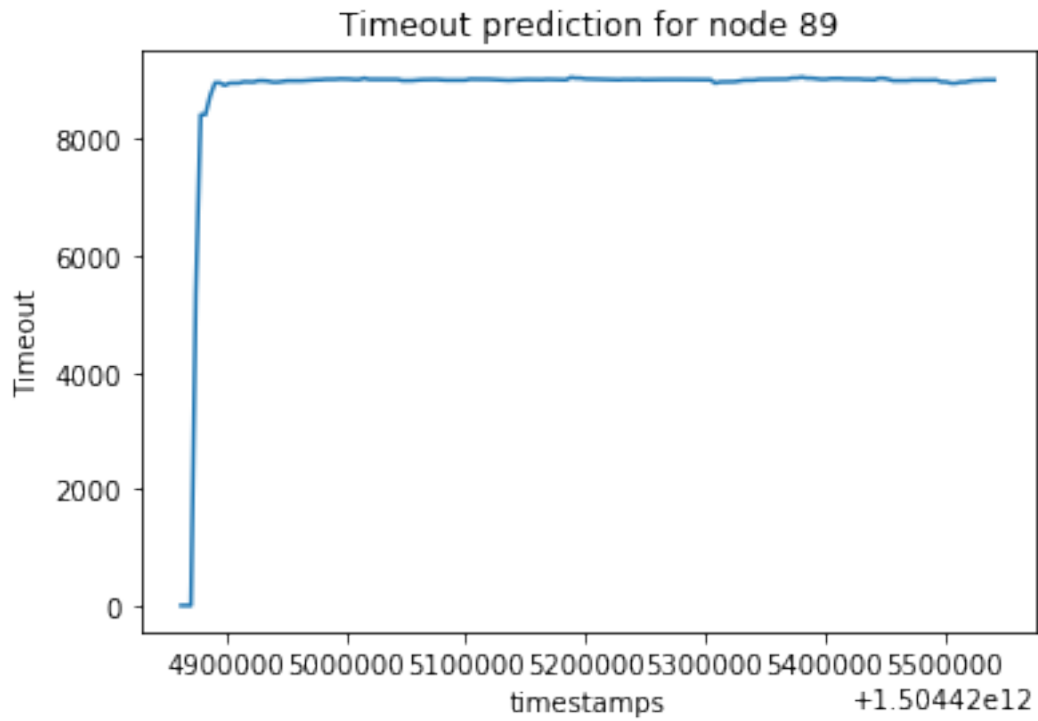


```
In [12]: rtt = rtt_data_with_mean.head().values
std = rtt_data_with_std.head().values
l = len(rtt)
width = 0.35
error_config = {'ecolor': '0.3'}
plt.bar(range(len(rtt)), rtt, width=width, color='r', yerr=std, error_kw=error_config,
plt.xlabel('GeoLocation')
plt.ylabel('Mean RTT')
plt.title('Mean RTT by GeoLocation with error-bars, only first 5 geo')
plt.legend()
plt.tight_layout()
```




1.4 Predictions for a single Node

```
In [13]: prediction_data = pd.read_csv('../data/backup/stats_mlfd/prediction_data.csv')
         node89 = prediction_data.loc[prediction_data['node'] == 89]
         time = node89["timestamp"].values
         prediction = node89["prediction"].values
         plt.plot(time, prediction)
         plt.ylabel('Timeout')
         plt.xlabel('timestamps')
         plt.title('Timeout prediction for node 89')
         plt.show()
```



1.5 RTT-data for a single Node

```
In [14]: node89 = rtt_data.loc[rtt_data['node'] == 89]
         time = node89["timestamp"].values
         prediction = node89["rtt"].values
         plt.plot(time, prediction)
         plt.ylabel('RTT')
         plt.xlabel('timestamps')
         plt.title('RTTs recorded for node 89')
         plt.show()
```




1.6 Analysis of number of suspected nodes with respect to time

```
In [15]: suspected_nodes_data = pd.read_csv('../data/backup/stats_mlfd/suspected_nodes.csv')
suspected_nodes_data.head()
```

```
Out[15]:
```

	timestamp	numberOfSuspectedNodes
0	1504424862212	0
1	1504424866440	0
2	1504424870471	61
3	1504424874490	20
4	1504424878510	0

```
In [16]: time = suspected_nodes_data["timestamp"].values
suspected = suspected_nodes_data["numberOfSuspectedNodes"].values
plt.plot(time, suspected)
plt.ylabel('number of suspected nodes')
plt.xlabel('timestamps')
plt.title('number of suspected nodes with respect to time')
plt.show()
```

```
In [17]: node_crashes = pd.read_csv('../data/backup/stats_mlfd/node_crashes.csv')
        node_suspicious = pd.read_csv('../data/backup/stats_mlfd/node_suspicious.csv')

In [18]: node_suspicious = pd.merge(node_suspicious, node_crashes, how="inner", on="node")
        node_suspicious = node_suspicious.sort_values('suspected', ascending=False).drop_duplicates()

In [19]: node_suspicious["detection_time"] = node_suspicious["suspected"] - node_suspicious["timestamp"]

In [20]: node_suspicious = node_suspicious[node_suspicious["detection_time"]>0] #Remove premature detections

In [21]: node_suspicious
```

```
Out[21]:
```

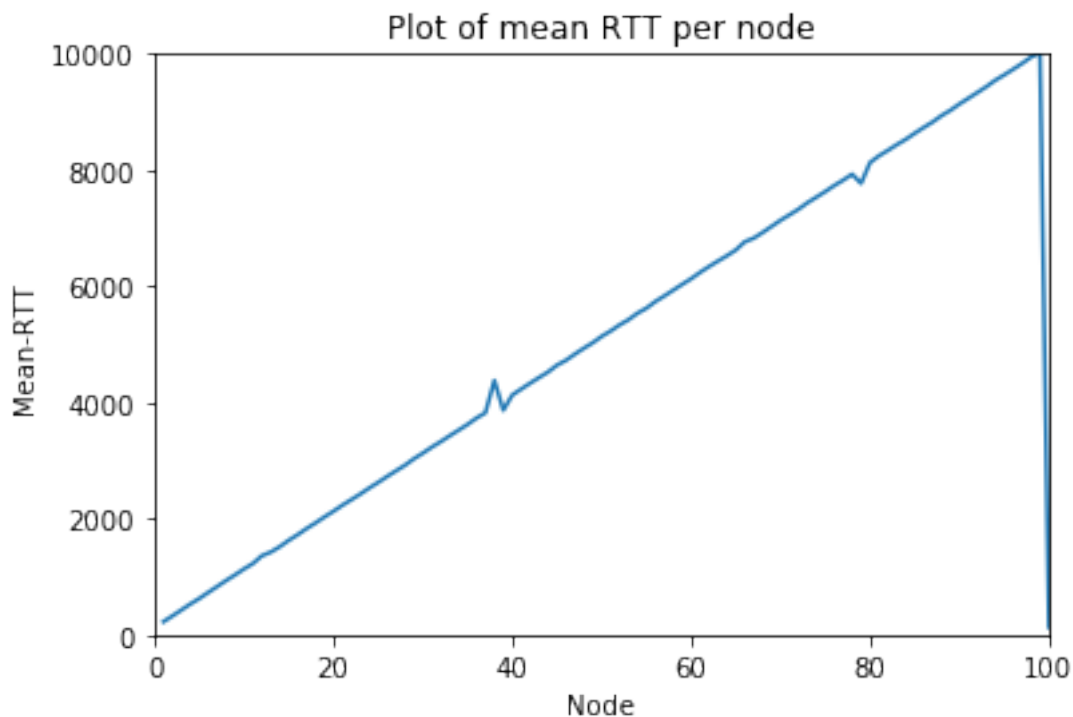
	node	suspected	timestamp	detection_time
1	91	1504425188050	1504425184031	4019
3	54	1504425119710	1504425115691	4019
6	66	1504424999110	1504424995091	4019
8	97	1504425453370	1504425441311	12059
10	73	1504424999110	1504424991072	8038
12	63	1504425401110	1504425393071	8039
14	74	1504425489550	1504425481511	8039
15	12	1504424894591	1504424890574	4017
16	14	1504425316690	1504425312671	4019
17	22	1504425401110	1504425397092	4018
18	20	1504425477490	1504425473471	4019

1.7 Mean RTT per node

```
In [22]: #rtt_data calculate mean per node and merge
        mean_rtt = rtt_data.groupby(['node', 'rtt'], as_index=False).mean().groupby('node')['rtt']

In [23]: matrix = mean_rtt.as_matrix
        mean = np.array(mean_rtt)

In [24]: plt.plot(mean_rtt)
        plt.axis([0, 100, 0, 10000])
        plt.ylabel('Mean-RTT')
        plt.xlabel('Node')
        plt.title('Plot of mean RTT per node')
        plt.show()
```




1.8 False suspicions per timeout

```
In [25]: node_suspensions2 = pd.read_csv('../data/backup/stats_mlfd/node_suspensions.csv')
        keys = ['suspected', "node"]
        i1 = node_suspensions2.set_index(keys).index
        i2 = node_suspensions.set_index(keys).index
        false_suspensions = node_suspensions2[~i1.isin(i2)]
        false_suspensions = false_suspensions.groupby("suspected").count()
        false_suspensions
```

```
Out [25]:
```

	node
suspected	
1504424866440	61

```
In [27]: mlfd_false_suspicious = false_suspicious.values
l = len(mlfd_false_suspicious)
width = 0.35
error_config = {'ecolor': '0.3'}
plt.bar(range(len(mlfd_false_suspicious)), mlfd_false_suspicious, width=width, color='b')
plt.xlabel('Time')
plt.ylabel('Number of false suspicions')
plt.title('Number of false suspicions over time')
plt.legend()
plt.tight_layout()
```




1.9 Time a node was suspected - time node actually crashed, plotted with respect to mean RTT

```
In [28]: mean_n = np.zeros(len(node_suspicious["node"]))
j = 0
for i in range(0, len(mean)):
    if(i in node_suspicious["node"].values):
        #print(i)
```

```
mean_n[j] = mean[i]
j = j+1
```

```
In [29]: plt.plot(mean_n, node_suspicious["detection_time"])
plt.ylabel('Detection Time')
plt.xlabel('Average RTT')
plt.title('Detection Time of crashes with respect to average RTT')
plt.show()
```



1.10 Implementation of an Offline-model to analyze the data, using Keras+Tensorflow

```
In [30]: from keras.models import Sequential
from keras.layers import Dense
from sklearn.metrics import r2_score

mean_geo_data = rtt_data.drop(["node", "timestamp"], 1)
#mean_geo_data = mean_geo_data[["geoLocation", "rtt", "mean"]].apply(pd.to_numeric)
df_train, df_test = np.split(mean_geo_data.sample(frac=1), [int(.8*len(mean_geo_data))])
Y_train = df_train["rtt"].values
Y_test = df_test["rtt"].values
X_train = df_train.drop("rtt", 1).values
X_test = df_test.drop("rtt", 1).values
```

Using TensorFlow backend.

```
In [31]: mean_geo_data.head()
```

```
Out[31]:
```

	geoLocation	rtt	mean
0	0	270.0	270.0
1	1	338.0	338.0
2	2	548.0	548.0
3	3	599.0	599.0
4	4	648.0	648.0

```
In [32]: # create model
```

```
model = Sequential()
model.add(Dense(20, input_dim=2, init='uniform', activation='tanh'))
model.add(Dense(1, init='uniform', activation='linear'))
```

```
# Compile model
```

```
model.compile(loss='mse', optimizer='adam', metrics=['accuracy'])
```

```
# Fit the model
```

```
model.fit(X_train, Y_train, nb_epoch=1000, batch_size=10, verbose=0)
```

```
# Calculate predictions
```

```
PredTestSet = model.predict(X_train)
```

```
PredValSet = model.predict(X_test)
```

```
# Save predictions
```

```
np.savetxt("trainresults.csv", PredTestSet, delimiter=",")
```

```
np.savetxt("valresults.csv", PredValSet, delimiter=",")
```

```
/home/limmen/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:3: UserWarning: Update
This is separate from the ipykernel package so we can avoid doing imports until
/home/limmen/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:4: UserWarning: Update
after removing the cwd from sys.path.
/home/limmen/anaconda3/lib/python3.6/site-packages/keras/models.py:844: UserWarning: The `nb_epoch`
warnings.warn('The `nb_epoch` argument in `fit` '
```

```
In [33]: #Plot actual vs predition for training set
```

```
TestResults = np.genfromtxt("trainresults.csv", delimiter=",")
```


```
plt.plot(Y_train,TestResults,'ro')
```

```
#Compute R-Square value for training set
```


```
TestR2Value = r2_score(Y_train,TestResults)
```

```
print("Training Set R-Square=", TestR2Value)
```

```
Training Set R-Square= -7.45569161964e-09
```



```
In [34]: plt.scatter(Y_train, TestResults,color=["red", "blue"])  
plt.show()
```




1.11 Another offline model, using sklearn

```
In [35]: from sklearn.model_selection import cross_val_predict
         from sklearn import linear_model

         lr = linear_model.LinearRegression()

         predicted = cross_val_predict(lr, X_train, Y_train, cv=10)

In [36]: fig, ax = plt.subplots()
         ax.scatter(Y_train, predicted, edgecolors=(0, 0, 0))
         ax.plot([Y_train.min(), Y_train.max()], [Y_train.min(), Y_train.max()], 'k--', lw=4)
         ax.set_xlabel('Measured')
         ax.set_ylabel('Predicted')
         plt.show()
```



1.12 Comparison MLFD vs EPFD

```
In [37]: node_crashes_epfd = pd.read_csv('../data/backup/stats_epfd/node_crashes.csv')
         node_suspensions_epfd = pd.read_csv('../data/backup/stats_epfd/node_suspensions.csv')
         node_suspensions_epfd
         node_suspensions_epfd = pd.merge(node_suspensions_epfd, node_crashes_epfd, how="inner", on="node_id")
         node_suspensions_epfd = node_suspensions_epfd.sort_values('suspected', ascending=False)
         node_suspensions_epfd["detection_time"] = node_suspensions_epfd["suspected"] - node_suspensions_epfd["start_time"]
         node_suspensions = node_suspensions[node_suspensions["detection_time"]>0] #Remove premature detections
         node_suspensions_epfd
```


```
Out [37]:
```

	node	suspected	timestamp	detection_time
4	45	1504433437756	1504433427237	10519
11	79	1504433921665	1504433911146	10519
15	42	1504433479836	1504433469317	10519
22	78	1504433574516	1504433563996	10520
25	20	1504433490356	1504433479836	10520
27	4	1504433711275	1504433700756	10519
33	71	1504433490356	1504433479836	10520
41	92	1504433753356	1504433742836	10520
45	34	1504433183715	1504433178697	5018

```
In [38]: #rtt_data calculate mean per node and merge
mean_rtt_epfd = rtt_data.groupby(['node', 'rtt'], as_index=False).mean().groupby('node')
matrix_epfd = mean_rtt.as_matrix
mean_epfd = np.array(mean_rtt_epfd)
```

```
In [39]: mean_n_epfd = np.zeros(len(node_suspensions_epfd["node"]))
j = 0
for i in range(0, len(mean)):
    if(i in node_suspensions_epfd["node"].values):
        #print(i)
        mean_n_epfd[j] = mean_epfd[i]
        j = j+1
```

```
In [40]: plt.plot(mean_n_epfd, node_suspensions_epfd["detection_time"], label="EPFD")
plt.plot(mean_n, node_suspensions["detection_time"], label="MLFD")
plt.ylabel('Detection Time')
plt.xlabel('Average RTT')
plt.title('Detection Time of crashes with respect to average RTT')
plt.legend()
plt.show()
```

```
In [41]: node_suspensions2_epfd = pd.read_csv('../data/backup/stats_epfd/node_suspensions.csv')
keys_epfd = ['suspected', "node"]
i1_epfd = node_suspensions2_epfd.set_index(keys_epfd).index
i2_epfd = node_suspensions_epfd.set_index(keys_epfd).index
false_suspensions_epfd = node_suspensions2_epfd[~i1_epfd.isin(i2_epfd)]
false_suspensions_epfd = false_suspensions_epfd.groupby("suspected").count()
false_suspensions_epfd
```

```
Out[41]:
```

	node
suspected	
1504433156016	96
1504433159076	5
1504433161096	17
1504433163616	20
1504433166636	31
1504433170156	39
1504433174176	10
1504433178696	25
1504433183715	30
1504433189236	16
1504433195256	25
1504433201776	15
1504433208796	20
1504433216315	6

```


1504433224336    20
1504433241876    11
1504433260416     5
1504433269936     1

```

```

In [42]: epfd_false_suspicious = false_suspicious_epfd.values
l = len(epfd_false_suspicious)
width = 0.35
error_config = {'ecolor': '0.3'}
plt.bar(range(len(epfd_false_suspicious)), epfd_false_suspicious, width=width, color='b')
plt.xlabel('Time')
plt.ylabel('Number of false suspicions')
plt.title('Number of false suspicions over time')
plt.legend()
plt.tight_layout()

```




```

In [43]: bar_width = 0.45
index1 = np.arange(1)
index2 = np.arange(len(mlfd_false_suspicious))
plt.bar(index2, mlfd_false_suspicious, width=width, color='r', error_kw=error_config, l
plt.bar(index1 + bar_width, epfd_false_suspicious, width=width, color='b', error_kw=err
plt.xlabel('Time')
plt.ylabel('Number of false suspicions')
plt.title('Number of false suspicions over time')

```

```
plt.legend()
plt.xticks(index1 + bar_width, ('MLFD', 'EPFD'))
plt.tight_layout()
```



In []: